

Milestone 4 – MLflow, FastAPI, Dashboard, Monitoring, and Local Deployment Documentation

1. MLflow Tracking Documentation

Overview

This module integrates **MLflow** to track machine learning experiments such as:

- Hyperparameters
- Evaluation metrics
- Artifacts (plots, models, etc.)
- Model versioning

The tracking is performed inside `train_with_mlflow_tracking()` during model training.

MLflow Experiment Setup

```
mlflow.set_experiment("Customer_Churn_Prediction")
```

Creates or selects an MLflow experiment named **Customer_Churn_Prediction**.

Function: `train_with_mlflow_tracking()`

Purpose

Train a RandomForestClassifier while automatically logging:

- Best hyperparameters
- Metrics (accuracy, precision, recall, F1, ROC-AUC)
- Trained model
- Feature importance plot

Workflow

1. **Start MLflow Run**
 2. `with mlflow.start_run(run_name="RandomForest_Churn_Prediction"):`
 3. **Perform GridSearchCV** for hyperparameter tuning.
 4. **Train best model** based on F1 score.
 5. **Log parameters**
 6. `mlflow.log_params(grid_search.best_params_)`
 7. **Log all important metrics**
 8. **Log model**
 9. `mlflow.sklearn.log_model(best_model, "random_forest_model")`
 10. **Log artifacts**
 - o Feature importance plot saved locally
 - o Logged as artifact to MLflow
-

Outcome

After running the training function, MLflow UI will show:

```
mlflow ui
```

Open: <http://127.0.0.1:5000>

2. FastAPI Service Documentation

The API exposes a production-ready machine learning inference service for churn prediction.

Project Structure

<code>main.py</code>	→ FastAPI application
<code>run_api.py</code>	→ Starts the API server
<code>encoders.pkl</code>	→ Saved encoders for categorical preprocessing
<code>customer_churn_model.pkl</code>	→ Serialized RandomForest model

FastAPI Application Overview

Endpoints

Endpoint	Method	Purpose
/	GET	Root endpoint, returns basic metadata
/health	GET	API and model health check
/model-info	GET	Returns feature names and model details
/predict	POST	Predict churn for one customer
/batch_predict	POST	Predict churn for multiple customers

Data Models

Input Model – CustomerData

Defines all required customer attributes such as:

- Demographics
- Account details
- Billing
- Service subscriptions

Output Model – PredictionResponse

Includes:

- churn_probability
 - churn_prediction (0 or 1)
 - prediction_class (“Churn” / “No Churn”)
 - confidence level
-

Preprocessing Explanation

The API performs:

- Label encoding using saved encoders
 - Handling unseen categories
 - Ensuring correct feature order
 - Filling missing features with defaults
-

Prediction Workflow

1. Receive JSON input
 2. Preprocess using `preprocess_input()`
 3. Predict class & probability
 4. Compute confidence level
 5. Return formatted response
-

3. Streamlit Dashboard Documentation

Overview

A real-time dashboard that interacts with the FastAPI backend to:

- Predict churn
 - Analyze CSV batch data
 - Display model info
 - Monitor API health
-

Dashboard Sections

1. Single Prediction

- Interactive form for entering customer data
 - Sends request to `/predict` endpoint
 - Displays results:
 - Probability bar plot
 - Risk level
 - Recommendations
-

2. Batch Analysis

Allows users to:

- Upload CSV
- Preview data
- (Future extension) send batch to API

3. Model Info

Displays:

- Model status
 - Loaded model type
 - Feature list
-

4. API Health

Shows:

- Whether API is running
 - Sample prediction test
 - Timestamps
-

Visualization

Uses matplotlib and seaborn for:

- Probability visualization
 - Customer risk breakdown
-

4. Model Monitoring System Documentation

1. ModelMonitor Class

Purpose:

Provides:

- Data drift detection
- Concept drift monitoring
- Logging predictions

- Generating monitoring reports
-

Core Methods

log_prediction()

Stores each inference:

- input data
- prediction
- probability
- optional actual label

Used for continuous model evaluation.

detect_data_drift()

Uses **Kolmogorov–Smirnov test** to compare:

- reference training data
- new incoming data

Flags drift when `p-value < 0.05`.

detect_concept_drift()

Compares model accuracy over last 100 predictions with baseline.

Drift is detected when accuracy drops more than 5%.

generate_monitoring_report()

Returns:

- avg probability
- distribution of risk levels
- drift detection summary
- total predictions

2. ModelRetrainingStrategy Class

Purpose:

Defines rules determining when the model should be retrained.

Triggers:

- Time-based schedule (default: every 30 days)
- Data drift detected
- Concept drift detected

`retrain_model()`

Retrains the model on new collected data.

5. Local Deployment Documentation

Install Dependencies

`pip install -r requirements.txt`

Start FastAPI

Run from root directory:

`uvicorn deployment.app:app --host 0.0.0.0 --port 8000 --reload`

FastAPI Docs:

→ <http://localhost:8000/docs>

Start Streamlit Dashboard

`streamlit run dashboard/app.py`

Dashboard URL:

→ <http://localhost:8501>